

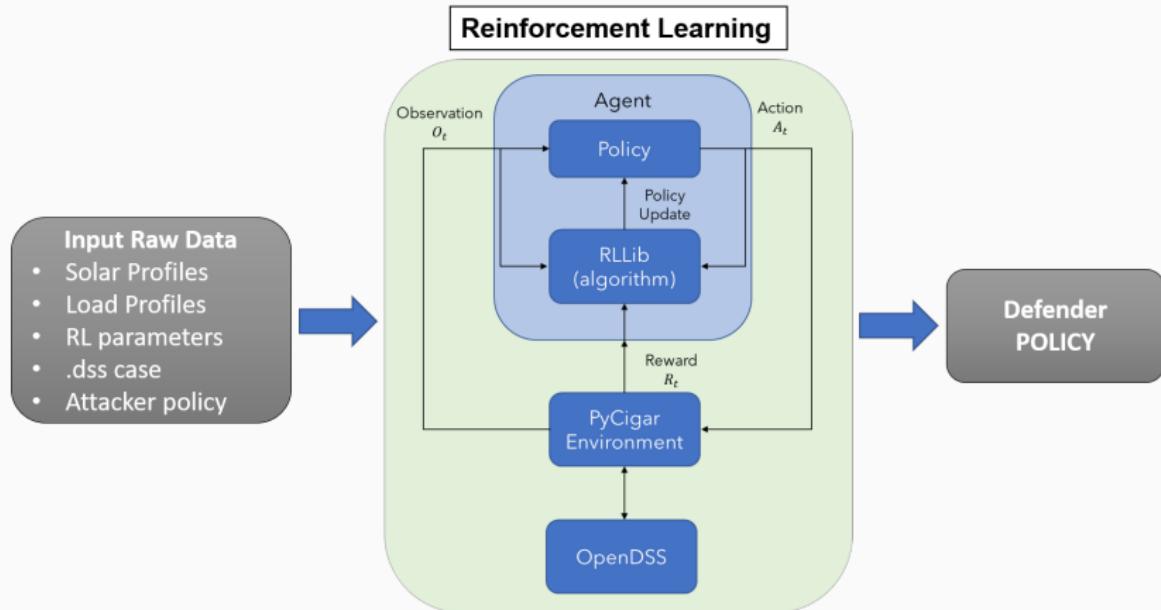
Synthetic Solar Data Generation and Linear Power Flow Solver for RL training

Ignacio Losada Carreño, Anna Scaglione

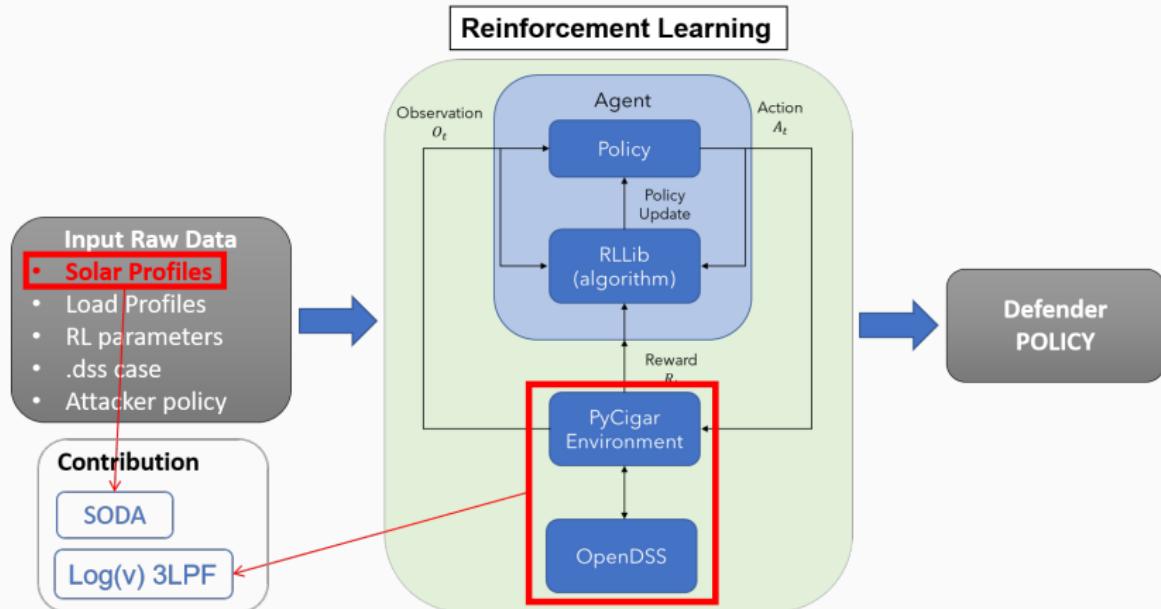
CIGAR Workshop. March 17, 2021

Arizona State University

Reinforcement learning for distribution systems



Improving RL training



SODA: An Irradiance-Based Synthetic Solar Data Generation Tool

Motivation and background

- RL agents trained on different solar and load conditions
- Low resolution datasets do not capture solar variability
- High resolution data not available (sparse locations)
- Combine a physics-based method (30-min resolution) with a stochastic model trained on PMU data to generate 1-second solar data

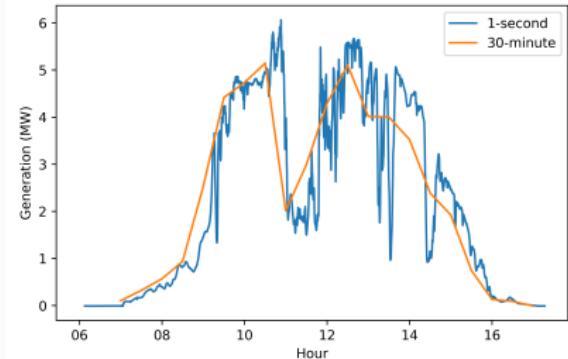


Figure 1: 1-second PMU data vs 30-min NSRDB data

⁰The National Solar Radiation Database (NSRDB) has a spatial resolution of 4km and a temporal resolution of 30 minutes

Motivation and background

Working hypothesis: The conditional distribution of solar power given the cloud density is the same for different locations

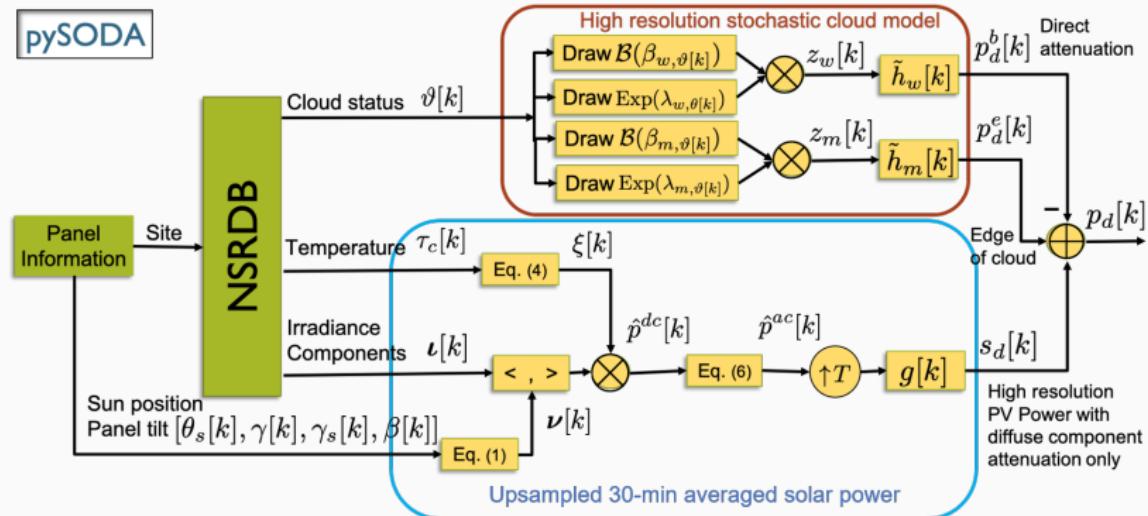


Figure 2: SODA block diagram

⁰<https://github.com/Ignacio-Losada/SoDa>

Showcasing SODA. Synthetic vs PMU data

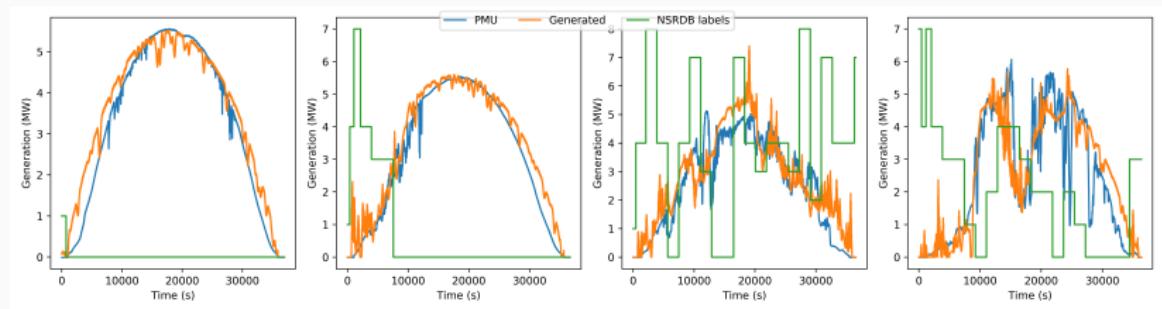


Figure 3: SODA vs 1-second PMU data from Riverside, CA

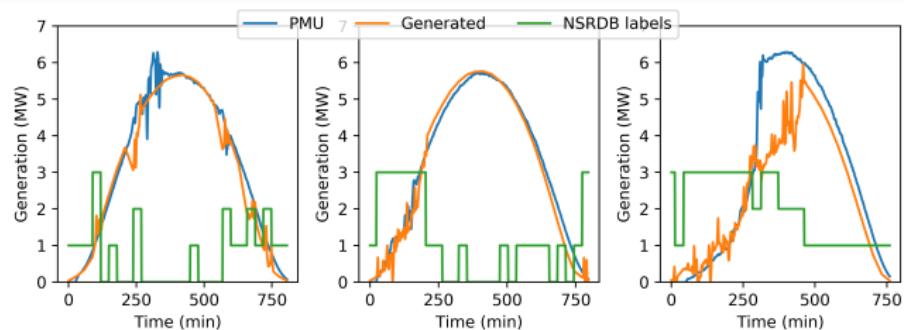


Figure 4: SODA vs 1-minute PMU data from Berkeley, CA

Footprint of SODA

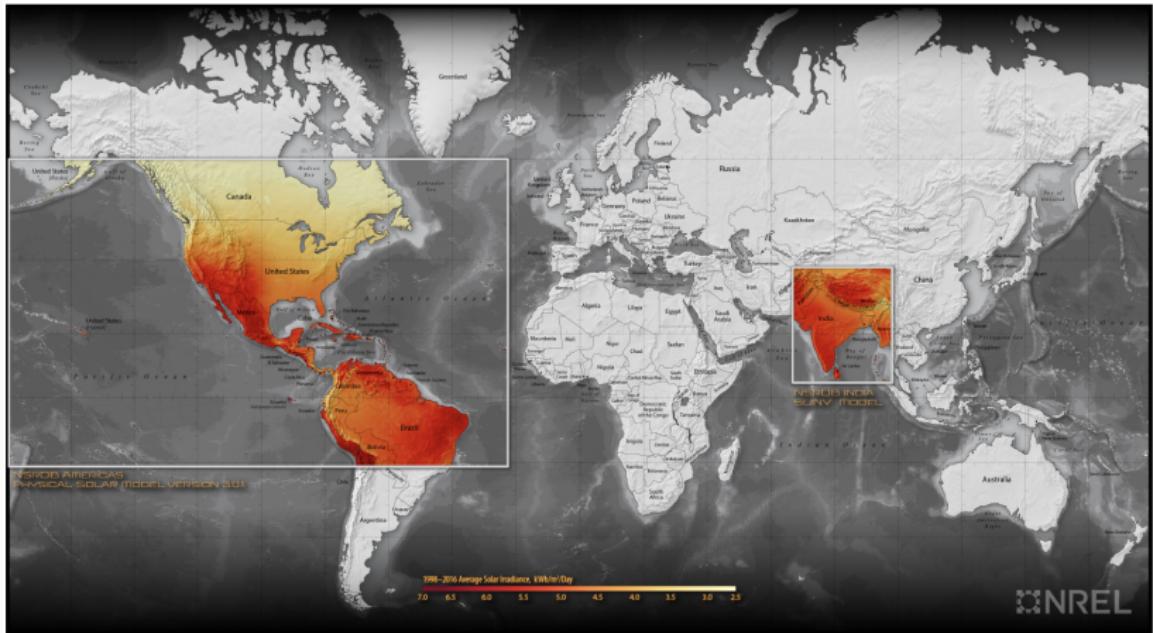


Figure 5: Area covered by the NSRDB

Limitations and summary

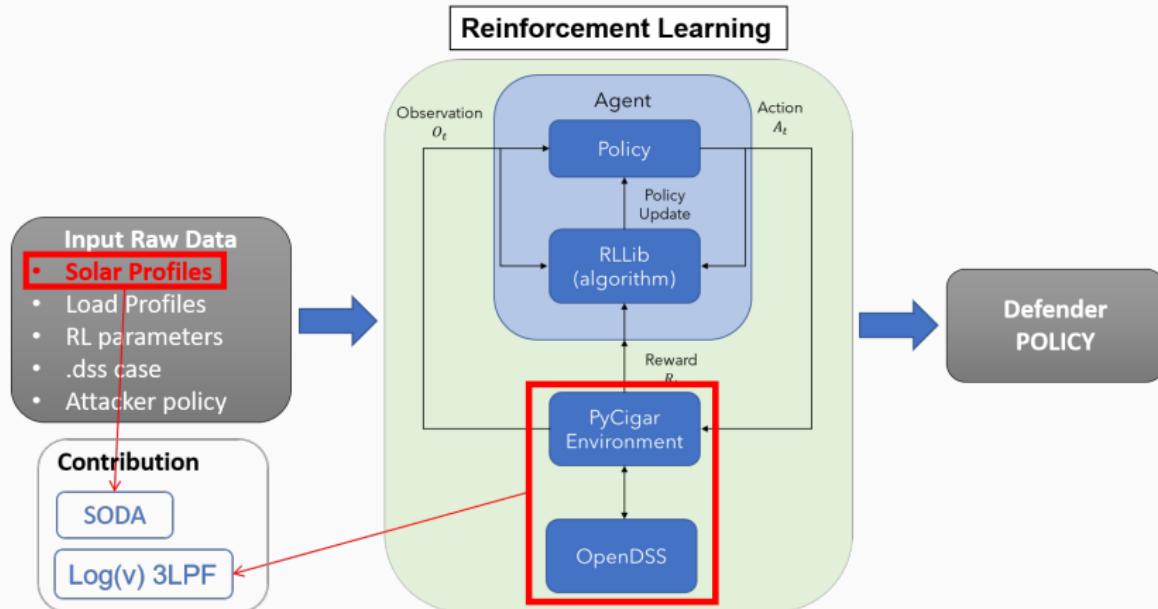
Limitations:

- Model does not account for time-of-day variability
- Mismatch NSRDB cloud labels
- Training was done with MW-scale data

Summary:

- We propose a stochastic model trained with PMU data that generates statistically representative 1-second resolution solar data
- This method, unlike NWP, scales for high resolutions
- <https://github.com/Ignacio-Losada/SoDa>

Improving RL training



Log(v) 3LPF: A linearized solution to train reinforcement learning algorithms for distribution systems

Reinforcement learning on distribution systems, PPO-Clip

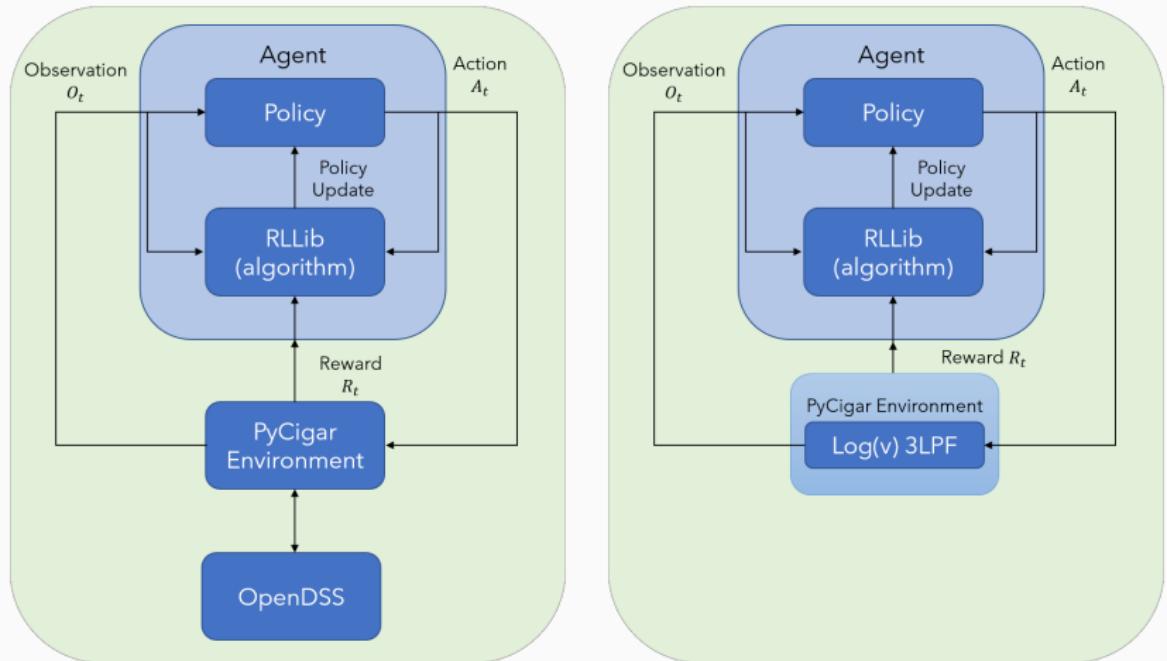


Figure 6: Current PyCigar modeling diagram (left) vs proposed architecture (right)

Reinforcement learning on distribution systems, PPO-Clip

- Power Flow (PF) equations are used to compute the rewards
- Rewards computed every training iteration, every time step and every action sampled by the RL algorithm
- Linear 3-phase unbalanced PF solver to speed up the training process

Formulation:

$$\text{OpenDSS} \rightarrow \mathbf{i} = \mathbf{Y}_{\text{bus}} \mathbf{v}$$
$$\text{Log(v) 3LPF:} \rightarrow \mathbf{s} = D(\mathbf{v} \mathbf{v}^H \mathbf{Y}_{\text{bus}}^H)$$

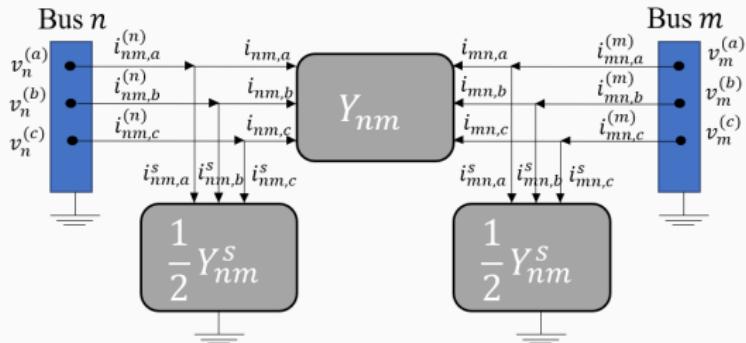


Figure 7: Pi-Model representation

Our contribution

- **Modeling capabilities**

- Embedded linear power flow solver
- ZIP load models
- Transformers, 3-phase and 240/120V center-tapped
- Voltage regulators and controls
- Capacitor banks and controls

- **Modeling accuracy**

- PF solution validated against OpenDSS
- Modeling of devices matches that of OpenDSS

Log(v) 3LPF

Kirchoff + Ohm + Losses:

$$\boxed{\mathbf{S}_{nm}^{(n)} = \mathbf{v}_n \mathbf{v}_n^H \left(\mathbf{Y}_{nm}^{(n)} + \frac{1}{2} \mathbf{Y}_{nm}^s \right)^H + \mathbf{v}_n \mathbf{v}_m^H \left(\mathbf{Y}_{nm}^{(m)} \right)^H}$$

We want to remove the non-linearity $\mathbf{v}_n \mathbf{v}_n^H$ from the equation that relates power flows to voltage

$$\mathbf{v}_n := [|v_n^a| e^{j\theta_a}, |v_n^b| e^{j\theta_b}, |v_n^c| e^{j\theta_c}]^T \quad (1)$$

$$|v_n^p| = e^{\log |v_n^p|}, \quad u_n^p := \log |v_n^p| \quad (2)$$

$$\boxed{v_n^p = e^{u_n^p} e^{j\theta_p}} \quad (3)$$

First-order Taylor expansion around 1 p.u. and 0 degree angles

Log(v) 3LPF

Non-linear ACPF: $S_{nm}^{(n)} = \mathbf{v}_n \mathbf{v}_n^H \left(Y_{nm}^{(n)} + \frac{1}{2} Y_{nm}^s \right)^H + \mathbf{v}_n \mathbf{v}_m^H \left(Y_{nm}^{(m)} \right)^H$

Expanding, dropping high-order terms, we may obtain

Log(v) 3LPF (Linear):

$$\tilde{s}_{nm} \approx \tilde{Y}_{\text{bus}} \mathbf{x} \quad \text{where} \quad \mathbf{x} \triangleq \begin{bmatrix} \mathbf{u} \\ \tilde{\theta} \end{bmatrix} \quad (4)$$

$$\mathbf{x} \approx \tilde{Y}_{\text{bus}}^{-1} \tilde{s}_{nm}$$

and we may recover the voltage phasors as follows

$$\mathbf{v} \approx \Delta_3 \text{diag}(e^{\mathbf{u}}) e^{j\tilde{\theta}_n}$$

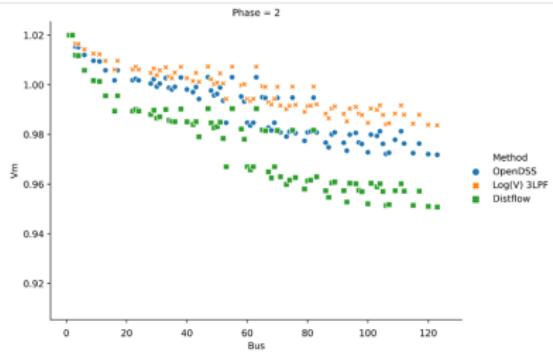
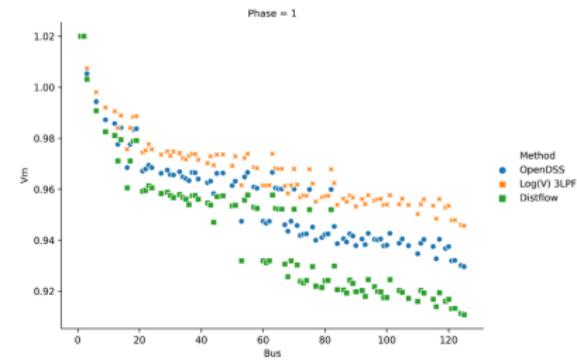
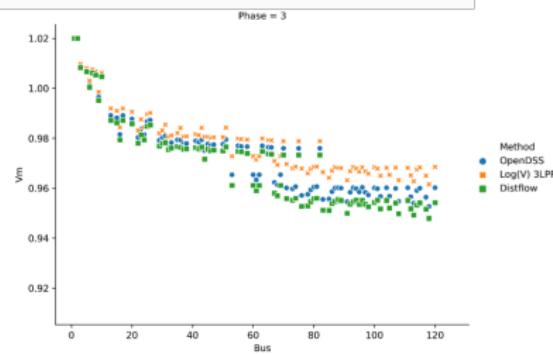
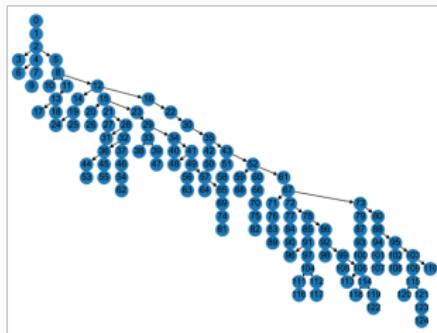
Outperforming OpenDSS when computing the PF solution

- Matrix Inversion, $\tilde{Y}_{\text{bus}}^{-1}$
 - Gauss-Jordan: $\mathcal{O}(n^3)$
 - LU decomposition (OpenDSS): $\mathcal{O}(\frac{2}{3}n^3)$
 - **Sherman-Morrison-Woodbury:** $\rightarrow \mathcal{O}((2n+1)k^2 + 3k^3)$
$$(A + UBV)^{-1} = A^{-1} - A^{-1}U(I + BVA^{-1}U)^{-1}BVA^{-1} \quad (5)$$
- Fixed-point iterations, $\mathcal{O}(n^2)$

Table 1: LU vs SMW in MFLOPS

	IEEE 13	IEEE 37	IEEE 123	IEEE 8500
LU	0.4	8.5	114	3311300
SMW	0.1	0.5	9.6	19.7
Ratio	4	17	12	168000

Solving one snapshot of the IEEE-123 test case



OpenDSS vs Log(V) 3LPF for testing policies

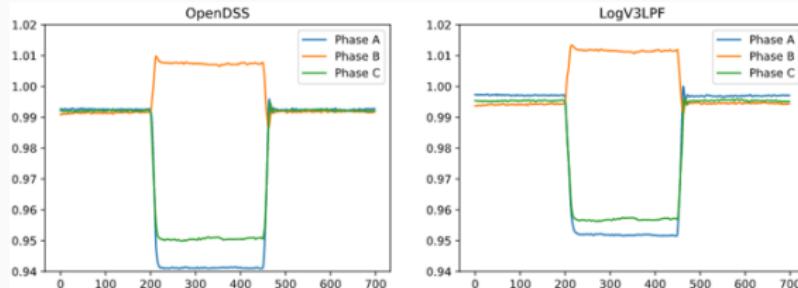


Figure 8: Voltage Imbalance attack

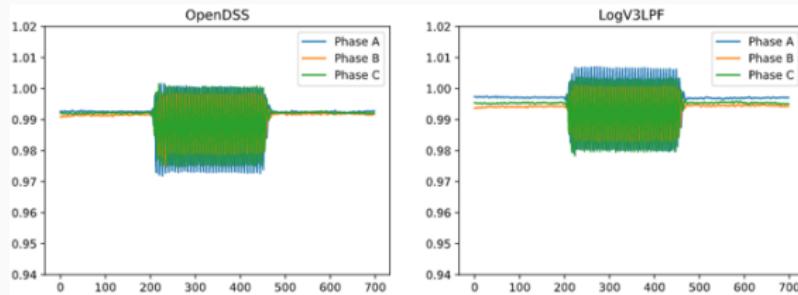
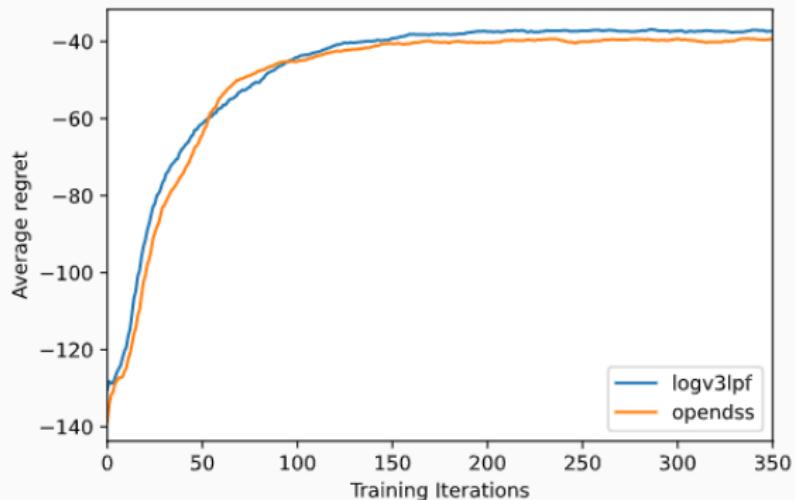
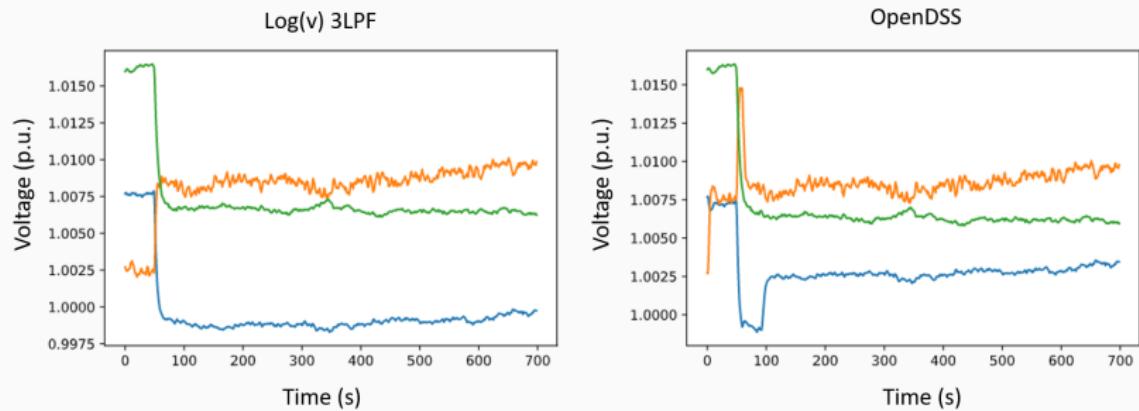


Figure 9: Voltage Oscillation attack

Training with OpenDSS vs Log(V) 3LPF. Average regret



Training with OpenDSS vs Log(V) 3LPF. Policy evaluation



Limitations and summary

Limitations:

- OpenDSS rarely inverts the system matrix
- OpenDSS convergence achieved in 3-4 iterations

Summary:

- We provide a linear PF solver with all necessary modeling capabilities
- Fast inverse computation
- Convergence is guaranteed (single iteration, linear)
- Model is, unlike OpenDSS, not sensitive to fast changes in boundary conditions